

```
In [2]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
df=pd.read_csv(r"C:\Users\Y.Saranya\Downloads\fiat500_VehicleSelection_Dataset.csv")
print(df)
```

	ID	model	engine_power	age_in_days	km	previous_owners	\
0	1	lounge	51	882	25000	1	
1	2	pop	51	1186	32500	1	
2	3	sport	74	4658	142228	1	
3	4	lounge	51	2739	160000	1	
4	5	pop	73	3074	106880	1	
...
1533	1534	sport	51	3712	115280	1	
1534	1535	lounge	74	3835	112000	1	
1535	1536	pop	51	2223	60457	1	
1536	1537	lounge	51	2557	80750	1	
1537	1538	pop	51	1766	54276	1	

	lat	lon	price
0	44.907242	8.611560	8900
1	45.666359	12.241890	8800
2	45.503300	11.417840	4200
3	40.633171	17.634609	6000
4	41.903221	12.495650	5700
...
1533	45.069679	7.704920	5200
1534	45.845692	8.666870	4600
1535	45.481541	9.413480	7500
1536	45.000702	7.682270	5990
1537	40.323410	17.568270	7900

[1538 rows x 9 columns]

```
In [3]: from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn import preprocessing, svm
df=df[['km', 'price']]
df.columns=['Km', 'Price']
df.head(10)
```

Out[3]:

	Km	Price
0	25000	8900
1	32500	8800
2	142228	4200
3	160000	6000
4	106880	5700
5	70225	7900
6	11600	10750
7	49076	9190
8	76000	5600
9	89000	6000

```
In [5]: df.tail()
```

Out[5]:

	Km	Price
1533	115280	5200
1534	112000	4600
1535	60457	7500
1536	80750	5990
1537	54276	7900

```
In [6]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1538 entries, 0 to 1537
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype  
---  -
0    Km      1538 non-null     int64  
1    Price   1538 non-null     int64  
dtypes: int64(2)
memory usage: 24.2 KB
```

```
In [7]: df.describe()
```

Out[7]:

	Km	Price
count	1538.000000	1538.000000
mean	53396.011704	8576.003901
std	40046.830723	1939.958641
min	1232.000000	2500.000000
25%	20006.250000	7122.500000
50%	39031.000000	9000.000000
75%	79667.750000	10000.000000
max	235000.000000	11100.000000

```
In [8]: df.shape
```

Out[8]: (1538, 2)

```
In [9]: df.isnull().sum()
```


Out[9]: Km 0
Price 0
dtype: int64

```
In [10]: x=np.array(df['Km']).reshape(-1,1)  
y=np.array(df['Price']).reshape(-1,1)
```

```
In [11]: df.dropna(inplace=True)  
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.3)  
regr=LinearRegression()  
regr.fit(X_train,y_train)  
regr.fit(X_train,y_train)  
print(regr.score(X_test,y_test))
```


0.7157349071219881

```
In [12]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [15]: udf=df[:][:500]
sns.lmplot(x="Km",y="Price",data=udf,order=1,ci=None)
```

Out[15]: <seaborn.axisgrid.FacetGrid at 0x11ddf6e1d20>




```
In [16]: udf.fillna(method='ffill',inplace=True)
X=np.array(udf['Km']).reshape(-1,1)
y=np.array(udf['Price']).reshape(-1,1)
udf.dropna(inplace=True)
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3)
regr.fit(X_train,y_train)
```

Out[16]:

▼ LinearRegression

LinearRegression()


```
In [17]: y_pred=regr.predict(X_test)
plt.scatter(X_test,y_test,color='y')
plt.plot(X_test,y_pred,color='b')
plt.show()
```



```
In [18]: from sklearn.linear_model import Ridge,Lasso,RidgeCV,LassoCV
```

```
In [19]: plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[19]: <Axes: >



```
In [21]: from sklearn.preprocessing import StandardScaler
features=df.columns[0:2]
target=df.columns[-1]
X=df[features].values
y=df[target].values
X_train,X_test,y_train,y_test=train_test_split(X,y,test_size=0.3,random_state=17)
print("The dimension of X_train is {}".format(X_train.shape))
print("The dimension of X_test is {}".format(X_test.shape))
scaler=StandardScaler()
X_train=scaler.fit_transform(X_train)
X_test=scaler.transform(X_test)
```

The dimension of X_train is (1076, 2)

The dimension of X_test is (462, 2)

```
In [22]: #Linear regression model
regr=LinearRegression()
regr.fit(X_train,y_train)
actual=y_test #actual value
train_score_regr=regr.score(X_train,y_train)
test_score_regr=regr.score(X_test,y_test)
print("\nLinear model:\n")
print("The train score for Linear model is {}".format(train_score_regr))
print("The test score for Linear model is {}".format(test_score_regr))
```

Linear model:

The train score for Linear model is 1.0

The test score for Linear model is 1.0

```
In [23]: #ridge regression model
ridgeReg=Ridge(alpha=10)
ridgeReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_ridge=ridgeReg.score(X_train,y_train)
test_score_ridge=ridgeReg.score(X_test,y_test)
print("\nRidge model:\n")
print("The train score for ridge model is {}".format(train_score_ridge))
print("The test score for ridge model is {}".format(test_score_ridge))
```

Ridge model:

The train score for ridge model is 0.9997095924476731

The test score for ridge model is 0.9997198323998524


```
In [24]: #using the linear cv model for ridge regression  
from sklearn.linear_model import RidgeCV  
#ridge cross validation  
ridge_cv=RidgeCV(alphas=[0.0001,0.001,0.01,0.1,1,10]).fit(X_train,y_train)  
#score  
print(ridge_cv.score(X_train,y_train))  
print(ridge_cv.score(X_test,y_test))
```

0.9999999999999676

0.9999999999999686


```
In [25]: #using the linear cv model for lasso regression  
from sklearn.linear_model import LassoCV  
#lasso cross validation  
lasso_cv=LassoCV(alphas=[0.0001,0.001,0.01,0.1,1,10],random_state=0).fit(X_train,y.  
#score  
print(lasso_cv.score(X_train,y_train))  
print(lasso_cv.score(X_test,y_test))
```

0.9999999877496772

0.9999999874481674

```
In [27]: #ridge regression  
plt.figure(figsize=(10,10))  
sns.heatmap(df.corr(),annot=True)
```

Out[27]: <Axes: >



```
In [28]: #lasso regression model
lassoReg=Lasso(alpha=10)
lassoReg.fit(X_train,y_train)
#train and test score for ridge regression
train_score_lasso=lassoReg.score(X_train,y_train)
test_score_lasso=lassoReg.score(X_test,y_test)
print("\nLasso model:\n")
print("The train score for lasso model is {}".format(train_score_lasso))
print("The test score for lasso model is {}".format(test_score_lasso))
```

Lasso model:

The train score for lasso model is 0.9999728562194999

The test score for lasso model is 0.9999728508562553

```
In [32]: pd.Series(lassoReg.coef_,features).sort_values(ascending=True).plot(kind="bar")
```

Out[32]: <Axes: >

